

Project 4: Reinforcement Learning

Train a Smartcab How to Drive

Reinforcement Learning Project for Udacity's Machine Learning Nanodegree.

Install

This project requires **Python 2.7** with the pygame library installed.

Code

Template code is provided in the `smartcab/agent.py` python file. Additional supporting python code can be found in `smartcab/enviroment.py`, `smartcab/planner.py`, and `smartcab/simulator.py`. Supporting images for the graphical user interface can be found in the `images` folder. While some code has already been implemented to get you started, you will need to implement additional functionality for the `LearningAgent` class in `agent.py` when requested to successfully complete the project.

Run

In a terminal or command window, navigate to the top-level project directory `smartcab/` (that contains this README) and run one of the following commands:

```
python smartcab/agent.py
python -m smartcab.agent
```

This will run the `agent.py` file and execute your agent code.

Implement a Basic Driving Agent

To begin, your only task is to get the smartcab to move around in the environment. At this point, you will not be concerned with any sort of optimal driving policy. Note that the driving agent is given the following information at each intersection:

The next waypoint location relative to its current location and heading. The state of the traffic light at the intersection and the presence of oncoming vehicles from other directions. The current time left

from the allotted deadline. To complete this task, simply have your driving agent choose a random action from the set of possible actions (None, 'forward', 'left', 'right') at each intersection, disregarding the input information above. Set the simulation deadline enforcement, `enforce_deadline` to False and observe how it performs.

QUESTION: Observe what you see with the agent's behavior as it takes random actions. Does the smartcab eventually make it to the destination? Are there any other interesting observations to note?

I have decided to implement a basic driving agent. It learns over time which actions lead to a positive reward and do not take any actions that have been observed to give negative rewards:

```
# TODO: Learn policy based on state, action, reward
if reward >= 0:
    self.valid_actions[self.state] = list(set(self.valid_actions.get(self.state, []) + [action]
    self.posreward += 1.0
else:
    self.negreward += 1.0
```

and

```
action = self.policy(self.state)
```

To test the basic agent, I have executed it 100 times (trials). In terms of performance metric, I use the the speed with which the cab can deliver its passenger to its destination as well as the ration of actions that give negative and positive rewards.

The results are not very encouraging: most of the times the agent does not reach the final destination. The statistics: - Overall, 2% of the actions gave a negative reward - Overall, 98% of the actions gave a positive reward - Only reached the destination in 2% of the trials

This means that the smartcab takes a defensive drive, but at the cost of not getting to the final destination often.

I have decided to compare this results with an even more naive approach: at any given point, the smartcab would take one of the any valid inputs.

```
action = self.policy(self.state, True)
```

The statistics for this version are not very encouraging: - Overall, 55% of the actions gave a negative reward - Overall, 45% of the actions gave a positive reward - Only reached the destination in 19% of the trials

This means that the smartcab needs more time to reach the final destination. However, a concerning factor is the fact that it takes risky actions, which is undesirable in this domain (e.g., accidents may occur)

Inform the Driving Agent

QUESTION: What states have you identified that are appropriate for modeling the smartcab and environment? Why do you believe each of these states to be appropriate for this problem? Please discuss the included states and the omitted ones as well.

Having in mind that at each intersection is a traffic light that either allows traffic in the North-South direction or the East-West direction and that the U.S. Right-of-Way rules apply

- On a green light, a left turn is permitted if there is no oncoming traffic making a right turn or coming straight through the intersection.
- On a red light, a right turn is permitted if no oncoming traffic is approaching from your left through the intersection.

I think that the following variables are necessary and sufficient to model state: `next_waypoint`, `light`, `oncoming`, and `left`. `Right` is not necessary because it is always possible to turn right, except if there is `oncoming` traffic from the `left` and the light is red, and that can be captured with the other state variables.

OPTIONAL: How many states in total exist for the smartcab in this environment? Does this number seem reasonable given that the goal of Q-Learning is to learn and make informed decisions about each state? Why or why not?

Given that there are 4 possibilities for `next_waypoint`, 2 possible values for `left`, `light`, and `oncoming`, there are 32 ($4 \times 2 \times 2 \times 2$) states. It seems reasonable, i.e., it is possible to learn the best actions for a given state in a reasonable amount of time, to learn the 32 possibilities and therefore make informed decisions.

Implement a Q-Learning Driving Agent

Your final task for this project is to enhance your driving agent so that, after sufficient training, the smartcab is able to reach the destination within the allotted time safely and efficiently. Parameters in the Q-Learning algorithm, such as the learning rate (α), the discount factor (γ) and the exploration rate (ϵ) all contribute to the driving agent's ability to learn the best action for each state. To improve on the success of your smartcab:

Set the number of trials, `n_trials`, in the simulation to 100. Run the simulation with the deadline enforcement `enforce_deadline` set to `True` (you will need to reduce the update delay `update_delay` and set the display to `False`). Observe the driving agent's learning and smartcab's success rate, particularly during the later trials. Adjust one or several of the above parameters and iterate this process. This task is complete once you have arrived at what you determine is the best combination of parameters required for your driving agent to learn successfully.

The Q-Learning policy implemented is best described using the following formula:

$$Q_{t+1}(s_t, a_t) \leftarrow \underbrace{Q_t(s_t, a_t)}_{\text{old value}} + \underbrace{\alpha_t(s_t, a_t)}_{\text{learning rate}} \cdot \left(\underbrace{R_{t+1}}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \underbrace{\max_a Q_t(s_{t+1}, a)}_{\text{estimate of optimal future value}} - \underbrace{Q_t(s_t, a_t)}_{\text{old value}} \right)$$

Figure 1: qlearning

Toggling between the simple Learning and the Q-Learning driving Agent is done by setting the following variable to `False/True` respectively:

```
qlearner = True
```

QUESTION: Report the different values for the parameters tuned in your basic implementation of Q-Learning. For which set of parameters does the agent perform best? How well does the final driving agent perform?

I tuned the Q-Learner to use the following parameters:

```
alpha = 0.7
gamma = 0.2
epsilon = 0.05
```

The statistics using these parameters are:

The statistics for this version are much better than the simple agent: - Overall, 5% of the actions gave a negative reward - Overall, 95% of the actions gave a positive reward - The smartcab reached the destination, within the time limit, in 97% of the trials. The budget it took to finish is plotted below:

I have also changed the learning parameter (alpha) to other values. As an example, below the statistics for `alpha = 0.5`:

The statistics for this version are much better than the simple agent: - Overall, 5% of the actions gave a negative reward - Overall, 95% of the actions gave a




Figure 2: boxplot




Figure 3: boxplot

positive reward - The smartcab reached the destination, within the time limit, in 95% of the trials. The budget it took to finish is plotted below:

Hence, reducing the learning rate means that there is a slight decrease in the number of trials that reach the final destination.

QUESTION: Does your agent get close to finding an optimal policy, i.e. reach the destination in the minimum possible time, and not incur any penalties? How would you describe an optimal policy for this problem?

On average, the Q-Learning-based agent reaches the destination in 9 of 10 trials, whereas the naive, basic agent reaches the destination in 2 of 10 trials only.

It is hard to calculate the optimal policy in such a dynamic environment. Moves that seem optimal may not really be because of the unpredictable behavior of other cars. I would however define an optimal policy as being one that go to the next way point without incurring negative rewards (as they may lead to accidents).

The agent has learned this type of optimal policy, and no actions with negative rewards are performed after a few trials (note that the y-axis is in log-scale).

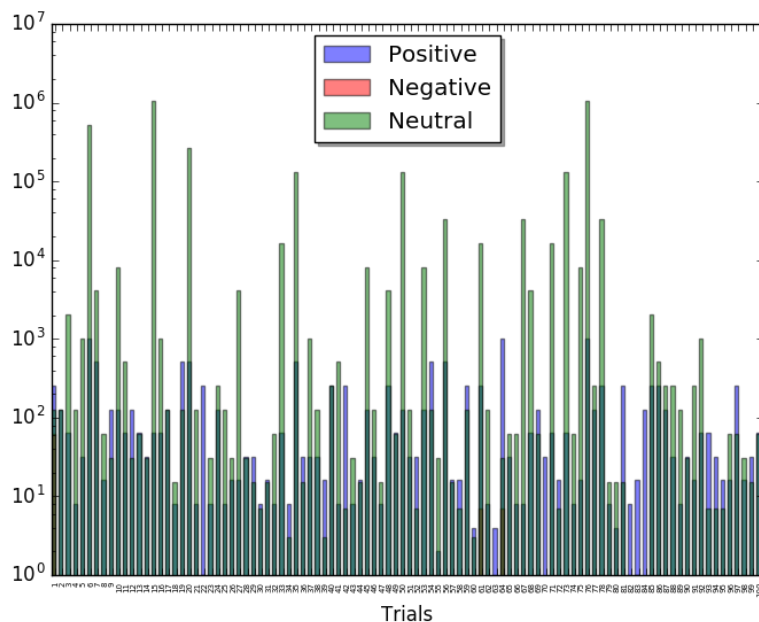


Figure 4: rewards